

Problem 1: Short Answer

(a) What is the running time of Dijkstras algorithm (with min-heap implementation) on a graph with n vertices and m edges?

Solution.

$$O((n + m) \log n)$$

(b) Dijkstra's algorithm is guaranteed to work correctly if the graph is directed with edge lengths that are non-negative (so they can be positive or zero, but not negative):

Solution.

TRUE

If there are negative length edges then $\text{dist}(v)$ can decrease after v has been explored, but if the lengths are non-negative then this cannot occur and hence Dijkstra's algorithm works correctly.

(c) Consider a MST T for undirected $G = (V, E)$. Now suppose for every edge e in G , we replace its edge weight $w(e)$ by $w(e) + 1$, so we increase every edge weight by $+1$. The tree T is guaranteed to still be a minimum spanning tree for this new weighted graph:

Solution.

TRUE

A spanning tree has exactly $n - 1$ edges, so increasing each edge weight by 1 increases the weight of all spanning trees by $n - 1$. It will not change the relative weights of two spanning trees, so a MST in the original graph will still be a MST in the new one.

(d) Consider a MST T for undirected $G = (V, E)$. Now suppose for every edge e in G , we replace its edge weight $w(e)$ by:

$$\hat{w}(e) = \begin{cases} 2w(e) & \text{if } w(e) > 100 \\ 0 & \text{if } w(e) \leq 100 \end{cases}$$

The tree T is guaranteed to still be a minimum spanning tree for this new weighted graph:

Solution.

TRUE

If an edge e was minimum across some cut (S, \overline{S}) in the original weighting scheme then it is still minimum across this same cut in the new weighting scheme (if it has weight ≤ 100 then there may be additional edges of the same new weight but e is still minimum).

Problem 2: MST

For an undirected graph $G = (V, E)$ with weights $w(e) > 0$ for each edge $e \in E$, you are given a MST T . Unfortunately one of the edges $e^* = (u, z)$ which is in the MST T is deleted from the graph G (no other edges change). Give an algorithm to build a MST for the new graph. Your algorithm should start from T .

Note: G is connected, and $G - e^*$ is also connected.

Explain your algorithm in words and use the algorithms from class as black-box subroutines. Clearly state and explain the running time of your algorithm. To receive credit, your algorithm should be correct and faster in $O()$ than building a MST from scratch (so don't use Prim's or Kruskal's, even part of Kruskal's).

Solution.

Let u and v be the two endpoints of e^* . Look at $T - e^*$. Let T_u and T_v be the two subtrees formed by removing e^* from T , such that $u \in T_u$ and $v \in T_v$. We can find the vertices in these 2 components by running DFS on $T - e^*$. We then go through all of the edges of the graph and take the minimum weight edge e' that has one endpoint in T_u and the other in T_v . We add e' to $T - e^*$ to get the new MST.

Running DFS takes $O(|V|)$ time since the graph $T - e^*$ has only $O(|V|)$ edges. Checking each edge takes $O(|E|)$ time. Therefore the algorithm takes $O(|V| + |E|)$ time.

Problem 3: Max-flow

Consider the following problem **Bipartite- k -Perfect**:

Input : Undirected bipartite graph $G = (V, E)$
and integer k where $1 \leq k \leq n$.

Output : A k -perfect subgraph if one exists, and NO if none exist in G .

Denote the 2 sides of the bipartite graph as $A = \{a_1, a_2, \dots, a_n\}$ and $B = \{b_1, \dots, b_n\}$; note $|A| = |B| = n$.

Explain how to reduce the **Bipartite- k -Perfect** problem to the max-flow problem. In other words, explain how to solve the bipartite k -perfect problem using the algorithm for solving max-flow as a black-box.

(a) Given an input $G = (V, E)$ and k to the Bipartite- k -Perfect problem, explain how you create the input to the max-flow problem (specify the graph G' and the edge capacities in this new graph). Do not do it for the above example, do it in general.

Solution.

To create G' , first add a source vertex s and a sink vertex t to G . Direct all edges of G to go from A to B and assign capacity 1 to each of these edges. Add edges of capacity k from s to each vertex in A , and from each vertex in B to t .

(b) Given a max-flow f^* for the flow network that you defined in part (a), explain how you use f^* to determine if G has a k -perfect subgraph.

Solution.

To check if G has a k -perfect subgraph, check if the max-flow in G' has size kn . If it does, then G has a k -perfect subgraph. The k -perfect subgraph is the set of edges between A and B with non-zero flow in f^* .

(c) What is the running time of your algorithm in terms of the original graph G where $n = |V|$ and $m = |E|$. State whether you are using the Ford-Fulkerson or Edmonds-Karp algorithm (only consider these two which we saw in class); faster is better.

Solution.

This algorithm uses Ford-Fulkerson. Constructing G' takes linear time ($O(n + m)$). Running Edmonds-Karp takes $O(nm^2)$, whereas running Ford-Fulkerson takes $O(mC)$ time. Since $C \leq kn$ then Ford-Fulkerson takes $O(knm)$ time which is better in this case.

Problem 4: Graph algorithms

You are given a directed graph $G = (V, E)$ where each edge e has a length $\ell(e) > 0$. Some of the vertices are marked special. Let $A = \{a_1, \dots, a_k\}$ be the subset of vertices marked special. The set A may be large, e.g., it may be that $|A| = O(n)$.

Devise an algorithm to compute the minimum total length from every vertex to the nearest special vertex. In other words, we want, for every $v \in V$ the minimum over all $a_i \in A$ of the distance from v to a_i .

Use the algorithms from class as a black-box subroutine. Get the best running time in $O()$ notation in terms of $n = |V|$ and $m = |E|$.

(a) There is **only 1 special vertex** so $|A| = 1$. Find the minimum distance from every vertex to this one special vertex a_1 .

Solution.

First, create the reverse graph $G^R = (V, E^R)$ ($E^R = \{(v, u) : (u, v) \in E\}$). Then, run Dijkstra's algorithm on G^R from a_1 . This finds the shortest path from a_1 to all other vertices. Reverse these paths to find the paths with minimum distance from every vertex to a_1 . The running time is $O((n + m) \log n)$ since 1 run of Dijkstra's algorithm is used.

(b) There are **many special vertices**.

Hint: transform the graph G so that you need just 1 run of Explore, DFS, BFS, or Dijkstra's algorithm.

Solution.

First, create the reverse graph $G^R = (V, E^R)$ ($E^R = \{(v, u) : (u, v) \in E\}$). Add a new vertex v to this graph, and add edges with 0 weight from v to each special vertex. Then, run Dijkstra's algorithm on this graph from v . Since the added edges have 0 weight, they do not affect the minimum distance paths. For a vertex z , its distance from v in this next graph will be the minimum distance from the closest special vertex. The running time is $O((n + m) \log n)$ since 1 run of Dijkstra's algorithm is used.